

Artificial Intelligence: Heuristic Search

Nathan Morgenstern, Seo Bo Shim

November 27, 2017

Contents

Introduction to GUI	2
Correct Solutions from instantiations of heuristic algorithm	3
Optimization of heuristic algorithm	4
Overall Optimizations	4
Heuristic Proposals	5
Best admissible/consistent heuristic	5
Inadmissible heuristics	6
Sub-optimality of Sequential A*	10

Introduction to GUI (a)

When the application first starts the user has the option to select either their own map, or a map that has been included by default. If the user wants to load their own map, they should select user generated map and input the file path into the file input text field. Otherwise, they can simply select a pre-generated map from the drop down menu along with its variation.

There are a total of 4 algorithms included in the main menu that the user can then select from to find a path from the start to goal. Options include: Uniform cost search, A*, Weighted A*, and Sequential A*. If the user selects weighted or sequential A* they should input two weights into the weight text fields. Once all the parameters are selected by the user, they can select confirm and the map along with its solution using the selected algorithm will be loaded.

Figure 1: The main menu of the application

Correct Solutions from instantiations of heuristic algorithm(b)

Figure 2: Uniform Cost Search example

Figure 3: A* Search example - Heuristic 1

Figure 4: Weighted A* Search example - Heuristic 2 $w_1 = 1.25$, $w_2 = 2$

Figure 5: Sequential A* Search example - $w_1 = 1.25$, $w_2 = 2$

Optimization of heuristic algorithm(c)

Overall Optimizations

When searching for the neighboring nodes to expand to, we can ignore the neighbor that is the parent because the parent was already explored. This saves the computation time of calculating the path and adding/removing from the binary heap.

The table that keeps the explored set of nodes is a 2 dimensional array. Since our grid is a 2d table, each node is assigned a coordinate. In the grid object that contains the graph of the map and all the nodes, we store a two dimensional array of nodes that represents the grid.

Also for keeping the fringe or the frontier of the graph search, we use a priority queue which is equivalent to a priority heap. We desire a minimal time for a search operation, and a binary heap provides this in $O(\log(n))$ time. Finding the highest element also takes $O(1)$, which is also significant in choosing this data structure to store our possible nodes in our frontier.

Since the overall algorithm is the same for all the types of searches, there were no optimizations specific to each type of search. Each search is different in how it calculates the value, which can consist of the actual path cost to the node or include the heuristic as well. IN the abstract algorithm we call an abstract method, where the implementation of this method is the only difference between the searches.

Heuristic Proposals(d)

Best admissible/consistent heuristic

The best heuristic can be achieved by considering the minimum path cost from the start to end goal.

We will consider the shortest path for three cases, and generalize the results

-A path from the start node to the goal node is a straight line.

The shortest possible path can occur if a river runs straight from the start to the goal node. This is possible because the river can possibly continue straight across the whole grid. This configuration allows the minimal past cost for this situation. The heuristic value of a node, the predicted cost from the node to the goal:

s is the current node, s_g is the goal node coordinates.

$t = 0.25$: multiplier for the type of block.

$d = 1.0$: multiplier for direction of movement.

$$h(s) = t * d * \max(|s^y - s_g^y|, |s^x - s_g^x|)$$

$$h(s) = 0.25 * 1 * \max(|s^y - s_g^y|, |s^x - s_g^x|)$$

-A path from the start node to goal node has diagonals and straight lines.

We note that it costs less to move non-diagonally on highways than than moving diagonally along non-rivers. The cost of moving diagonally is:

$$h(s) 4 * 0.25 * \sqrt{2} * \sqrt{2} * \min(|s^y - s_g^y|, |s^x - s_g^x|)$$

We note that moves can only be made vertically or horizontally on a highway. So we calculate the total

amount of vertical and horizontal moves.

$$h(s) = t * d * (vert_{dist} + horiz_{dist})$$

$$h(s) = 0.25 * 1 * (|s^y - s_g^y| + |s^x - s_g^x|)$$

We can clearly see that this heuristic is minimized.

Inadmissible heuristics

We argue that inadmissible heuristics should be used for the A* algorithm on this graph because the range of possible values for the path cost in this grid is high. Since the range is high, to select a admissible heuristic, we must shoot low and calculate the heuristic conservatively to not exceed the minimum path cost. Ideally, we wish for the minimum path cost to be close to the average, or the range to be smaller so that the heuristic value is not much smaller than the cost. Condition:

$$h(s) \leq c(s, s') + h(s')$$

$$h(s) - h(s') \leq \min(c(s, s'))$$

To find the maximal heuristic that fits this condition, we must find the minimum possible cost value. With a high range of values for the cost we find that:

$$h(s) - h(s') \ll \text{average}(c(s, s'))$$

In other words, the case where the path is strictly on a river or highway is very unlikely, yet it is still possible. This in turn reduces the maximum heuristic value to force the heuristic to be admissible, and in turn renders the heuristic value to minimally affect the estimated path cost. This is why the admissible and consistent A* implementation almost exactly follows the Uniform Cost Search algorithm.

H2: Type-based Heuristic

Currently, the abstract algorithm tends to follow rivers because the path cost to these rivers is small. However this is computationally expensive, as these nodes will have to continuously be expanded along and not along the river. We can argue that the path should follow the river as long as this does not increase the Euclidian distance from the current node to the goal.

If the current node is on a river, then the heuristic can be calculated to be:

$$h(s) = 0.38 * 1 * (|s^y - s_g^y| + |s^x - s_g^x|)$$

If the river takes us away from the goal, then the heuristic will be large enough that a step of 1 will be enough to have a lower path cost than following the river away from the goal.

The result is that nodes along the river going away from the goal will be expanded less, thus reducing computation time when the path is on a river.

When the current node is a hard to traverse cell, we make the assumption that this region will be full, so we avoid these cells by raising the heuristic value of. We calculate that traversing around a 31x31 grid of hard-to-traverse cells is slightly shorter than traversing through the 31x31 grid. Since there's only a 50 percent chance of a hard-to-traverse cell of being. Note: traveling through the difficult blocks on average will yield a path cost of 1.5.

Traversing around the 31x31 block: $30 + 30 + \sqrt{2} = 61.41$

Traversing through the 31x31 block

$$\sqrt{2} * 30 * (1 + (1 * (0.50))) = 65.76$$

$61.41 < 65.76$. It is advantageous to travel around the 31x31 grid. Though it may have a higher path cost to travel around the block, because with a larger number of blocks. So this heuristic should be designed to have the path best avoid these hard-to-traverse cells, not necessarily avoid them completely.

$$h(s) = 1.5 * 1 * (|s^y - s_g^y| + |s^x - s_g^x|)$$

When the current node is a regular unblocked cell:

$$h(s) = t1 * d1 * \max(\text{vert_dist}, \text{horiz_dist}) + t2 * d2 * \min(\text{vert_dist}, \text{horiz_dist})$$

$$h(s) = 1 * 1 * \min(|s^y - s_g^y|, |s^x - s_g^x|) + 1 * \sqrt{2} * \max(|s^y - s_g^y|, |s^x - s_g^x|)$$

s is the current node, s_g is the goal node coordinates.

$t1 = t2 = 1$: multiplier for the type of block.

$d1 = 1.0$, $d2 = \sqrt{2}$: multiplier for direction of movement.

This heuristic is advantageous because since it uses the current node type, it is computationally inexpensive.

H3: Expected Value

This method will take into account the number of blocked cells and the number of hard to traverse cells. The value that's generated from the heuristic will take the probability of each type of cell occurring and the required cost to traverse this path. The heuristic value will be the average cost of each step combined with a value that's proportional to the node's distance to the goal. The Euclidian distance will be used.

Average cost of step = summation of: probability of scenario * cost of scenario.

$$\text{average}(s) = \sum P(s) * C(s)$$

]

$$\text{average}(s) = (0.2 * 1.91 + 0.8 * (0.8 * 1 + 8 * 0.05 * 0.5 * 1.5)) = 0.3 + 0.8 = 1.482$$

For each blocked cell, we can consider the cost to navigate around it as the cost of having the blocked

cell. We consider the different cases of how a blocked cell can be overcome. Approaching from a corner:

$$(2 + \sqrt{2}) + 2 * (1 + \sqrt{2}) + 2 * 2 + 2 * 1 / 7 = 2.03$$

Approaching the blocked cell perpendicularly:

$$(2 * \sqrt{2}) + 2 * (1 + \sqrt{2}) + 2 * (\sqrt{2}) + 2 * 1 / 7 = 1.78$$

Averaging those two values together: 1.91 cost for each blocked cell.

There is an 80 percent chance that a path will be unblocked. For each 31x31 grid of 50 percent hard to traverse blocks, this covers up roughly 5 percent of the 120 x 160 grid. Eight 31 x 31 grids result in about 40 percent of the grid covered, consequently about 20 percent of the map is covered by hard to traverse cells.

Not exactly 20 percent of the map is covered by hard to traverse cells though, as there is a chance of overlap; each hard to traverse cell is not uniquely placed. There's also a chance the cell will land on the edge and not all of the 31x31 chosen cells will be on the grid.

This heuristic also ignores the probability of there being a river, so this heuristic will overestimate the path cost of reaching the goal. This heuristic can possibly be optimized by having a constant value that can be adjusted to account for the probability of rivers occurring in the optimal path.

H4: Manhattan Distance

This heuristic will take the Manhattan distance from the current node to the end goal. This heuristic is useful when diagonal steps are not possible in the grid. This heuristic is similar to the Euclidean distance, except we assume that diagonal steps are not taken.

This heuristic will focus on finding the straightest route from the node to the goal. This heuristic is expected to cause the algorithm to be quick.

H5: Euclidean Distance

This heuristic will take the Euclidean distance from the current node to the end node. In other words, the distance will be calculated as a line directly from the node to the goal node. This heuristic may provide the most accurate representation of the cost from the node to the goal disregarding the presence of rivers.

This heuristic is useful because it'll prioritize the actual distance of the goal from the current node. It will result in little nodes being expanded and focus on the straightest route to the goal.

Heuristic Results and Comparisons(e and f)

Experimental Results

Algorithm	Mean Time	Mean Nodes	Optimal Path Length
Uniform Cost	21.9412 ms	12907	179.66

Algorithm	Mean Time	Mean Nodes	Percent from optimal
A* - H1	21.6008 ms	10374	0.000 %
A* - H2	2.9807 ms	460	-5.755 %
A* - H3	0.9796 ms	217	-28.877 %
A* - H4	1.8209 ms	754	-16.932 %
A* - H5	3.3776 ms	1660	-15.886 %

Algorithm	Mean Time	Mean Nodes	Percent from optimal
W = 1.25 A* - H1	23.4999 ms	9486	-1.747 %
W = 1.25 A* - H2	4.2639 ms	417	-4.386 %
W = 1.25 A* - H3	1.9081 ms	136	-33.074 %
W = 1.25 A* - H4	1.4664 ms	288	-30.669 %
W = 1.25 A* - H5	0.9395 ms	376	-22.821 %

Algorithm	Mean Time	Mean Nodes	Percent from optimal
W = 2 A* - H1	14.2124 ms	6773	-5.143 %
W = 2 A* - H2	2.0119 ms	398	0.256 %
W = 2 A* - H3	0.4456 ms	120	-34.220%
W = 2 A* - H4	0.5921 ms	131	-31.882 %
W = 2 A* - H5	0.3917 ms	127	-33.819 %

Algorithm	Mean Time	Mean Nodes	Percent from optimal
Sequential A*	10.0905 ms	817	-29.967 %

Comparison

As expected, the consistent and admissible heuristic, H1, is expected to behave

Sequential A*

Implementation

The implementation of the sequential A* algorithm is efficient because it utilizes various data structures to minimize the time needed to access the data. The frontier priority queue and the 2D closed array for each heuristic are kept in a list. This ArrayList only holds five elements, equivalent to the number of different heuristics. The ArrayList allows for random access based on the index, which is useful because the algorithm switches between the different heuristics constantly so the access times should be minimized.

Again, a priority heap is utilized to minimize the search time and the access time for the minimal value in the heap.

The parents of all the nodes are only updated when the goal is reached. A 2D array is kept to temporarily store the parent of each node at a coordinate, and only at the end the node's actual parents are updated. Parents of a node can change often while the optimal cost path is in the process of being found, so this reduces the need to have access to the node object until the end.

Sub-optimality of Sequential A*

$$g_0(s) \leq w_1 * c^*(s)$$

In the anchor search of the sequential A*, we know that the heuristic is admissible and consistent. The cost from the start to a node is less than the optimum cost * the weight

$$Keys(s) \leq w_1 * g^*(s_{goal})$$

$$Keys(s) = g(s) + w_1 * h(s)$$

$$g(s) + w_1 * h(s) \leq w_1 * g^*(s_{goal})$$

$$h(s) \leq c(s', s) + h(s')$$

The anchor key's heuristic value is admissible and consistent, so it will never overestimate the path to the goal. Applying this condition to the starting node:

$$h(s_{goal}) \leq g(s_{goal}) + 0$$

We can apply this to the condition:

$$g(s) + w_1 * g(s) \leq w_1 * g^*(s_{goal})$$